

SIIM-ACR Pneumothorax Segmentation

1. Business Problem

1.1 Description

A pneumothorax or collapsed lungs is a medical condition that is responsible for making the people suddenly gasp for air, and feel helplessly breathless for no apparent reason. It can be a complete lung collapse or a collapse of a portion of the lung. It is usually diagnosed by a radiologist with several years of experience on a chest x-ray; which sometimes is very difficult to confirm. Our goal is to classify(if present segment) pneumothorax from a set of chest radiographic images.

This is a Kaggle problem. (<https://www.kaggle.com/c/siim-acr-pneumothorax-segmentation/overview>)

1.2 Problem Statement

Classify pneumothorax from a set of chest radiographic images and if present segment the regions of pneumothorax.

1.3 Sources

- Source : <https://www.kaggle.com/c/siim-acr-pneumothorax-segmentation/overview>
-

1.4 Real world/Business Objectives and Constraints

- Classify pneumothorax and if present segment it.
- Maximize the overlap between the actual mask and predicted mask(Dice).
- No such latency concerns.

2. Machine Learning Problem

2.1 Data

2.1.1 Data Overview

- Data is present in under folder SIIM. It contains three folders and files,

- dicom-images-test
- dicom-images-train
- train-rle.csv

Note : I downloaded the data from <https://www.kaggle.com/seesee/siim-train-test> as the data is removed from the Cloud Healthcare API.

2.2 Mapping the real world problem to a Deep Learning Problem

2.2.1 Type of Deep Learning Problem

It is a basically a semantic image segmentation problem, where we have to segment areas of pneumothorax.

2.2.2 Performance Metric

Source: <https://www.kaggle.com/c/siim-acr-pneumothorax-segmentation/data>

Metric:

- Focal Loss + Dice Loss
 - Focal loss is very good for imbalanced data as it focuses more on hard examples than easy examples. Our data is very imbalanced as the area having pneumothorax is very small. Dice is best metric overall. So we are going for the combination of the two.

2.2.3 Necessary Imports

```
import numpy as np

def mask2rle(img, width, height):
    rle = []
    lastColor = 0;
    currentPixel = 0;
    runStart = -1;
    runLength = 0;

    for x in range(width):
        for y in range(height):
            currentColor = img[x][y]
            if currentColor != lastColor:
                if currentColor == 255:
                    runStart = currentPixel;
                    runLength = 1;
                else:
                    rle.append(str(runStart));
                    rle.append(str(runLength));
                    runStart = -1;
                    runLength = 0;
                    currentPixel = 0;
            elif runStart > -1:
                runLength += 1
            lastColor = currentColor;
            currentPixel+=1;

    return " ".join(rle)

def rle2mask(rle, width, height):
    mask= np.zeros(width* height)
    array = np.asarray([int(x) for x in rle.split()])
    starts = array[0::2]
    lengths = array[1::2]

    current_position = 0
```

```

for index, start in enumerate(starts):
    current_position += start
    mask[current_position:current_position+lengths[index]] = 255
    current_position += lengths[index]

return mask.reshape(width, height)

```

```
!pip install pydicom
```

```

Collecting pydicom
  Downloading https://files.pythonhosted.org/packages/f4/15/df16546bc59bfca390cf072d4731
    |████████████████████████████████████████| 1.9MB 6.2MB/s
Installing collected packages: pydicom
Successfully installed pydicom-2.1.2

```

```

import pydicom
import numpy as np
import matplotlib.pyplot as plt
import os
import pandas as pd
import glob
from tqdm import tqdm
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
import gc
gc.collect()
import cv2
import tensorflow as tf
import keras
from PIL import Image, ImageDraw
from PIL import ImagePath
import imgaug.augmenters as iaa
from skimage import exposure

```

1. Creating train and test dataset

```

data_rle_train = pd.read_csv("./siim/train-rle.csv")
data_rle_train.head()

```

	ImageId	EncodedPixels
0	1.2.276.0.7230010.3.1.4.8323329.6904.151787520...	-1
1	1.2.276.0.7230010.3.1.4.8323329.13666.15178752...	557374 2 1015 8 1009 14 1002 20 997 26 990 32 ...
2	1.2.276.0.7230010.3.1.4.8323329.14666.15178752...	.

#The second column contains a space in front of it so manually rewriting it

```
data_rle_train.columns = ['ImageId', 'EncodedPixels']
```

<https://www.kaggle.com/jesperdramsch/intro-chest-xray-dicom-viz-u-nets-full-data>

```
def getInfoDICOM(path, toPrint = False, train = True):
```

```
    info = {}
```

```
    path = os.path.join(path)
```

```
    #reading the data using methods of pydicom
```

```
    data = pydicom.dcmread(path)
```

```
    info['path'] = path
```

```
    info['age'] = data.PatientAge
```

```
    info['sex'] = data.PatientSex
```

```
    info['ImageId'] = data.SOPInstanceUID
```

```
if train: #test doesn't have encoded pixels data, we have to predict them
```

```
    #SOPInstanceUID contains the storage type which resembles the image ids given in train
```

```
    encodedPixels = data_rle_train[data_rle_train['ImageId'] == data.SOPInstanceUID]["EncodedPixels"]
```

```
    info['encodedPixels'] = encodedPixels
```

```
    info['lenOfEncodedPixels'] = len(encodedPixels)
```

```
    #this is for visualization purpose
```

```
    if '-1' in encodedPixels or len(encodedPixels) == 0:
```

```
        info['has_pneumothorax'] = 0
```

```
    else:
```

```
        info['has_pneumothorax'] = 1
```

```
if toPrint:
```

```
    print("Path.....:", path)
```

```
    print("Patient's Name.....:", data.PatientName)
```

```
    print("Patient's Id.....:", data.PatientID)
```

```
    print("Patient's Age.....:", data.PatientAge)
```

```
    print("Patient's Sex.....:", data.PatientSex)
```

```
    print("Original X Ray")
```

```
    plt.figure(figsize=(10,10))
```

```
    plt.imshow(data.pixel_array, cmap=plt.cm.bone)
```

```
    plt.show()
```

```
return info
```

```
#creation of train dataset
#https://stackoverflow.com/questions/33747968/getting-file-list-using-glob-in-python
filesList = glob.glob('./siim/dicom-images-train/**/*.dcm')
train = pd.DataFrame()
train_list = []
for file in tqdm(filesList):
    info = getInfoDICOM(file, False, True)
    train_list.append(info)
train = pd.DataFrame(train_list)
```

```
100%|██████████| 12089/12089 [00:40<00:00, 296.59it/s]
```

```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12089 entries, 0 to 12088
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   path                   12089 non-null object
1   age                    12089 non-null object
2   sex                    12089 non-null object
3   ImageId                12089 non-null object
4   encodedPixels          12089 non-null object
5   lenOfEncodedPixels     12089 non-null int64
6   has_pneumothorax      12089 non-null int64
dtypes: int64(2), object(5)
memory usage: 661.2+ KB
```

```
train.head(5)
```

	path	age	sex	ImageId	enco
0	./siim/dicom-images-train/1.2.276.0.7230010.3....	73	M	1.2.276.0.7230010.3.1.4.8323329.6277.151787519...	
1	./siim/dicom-images-train/1.2.276.0.7230010.3....	10	F	1.2.276.0.7230010.3.1.4.8323329.7035.151787520...	
2	./siim/dicom-images-train/1.2.276.0.7230010.3....	64	M	1.2.276.0.7230010.3.1.4.8323329.14139.15178752...	[978: 20 ! 33 !
3	./siim/dicom-images-train/1.2.276.0.7230010.3....	22	F	1.2.276.0.7230010.3.1.4.8323329.3216.151787517...	

[978:
20 !
33 !

[256:

```
#creation of test dataset
filesList = glob.glob('./siim/dicom-images-test/**/*.dcm')
test = pd.DataFrame()
test_list = []
```

```

for file in tqdm(filesList):
    info = getInfoDICOM(file, False, False)
    test_list.append(info)
test = pd.DataFrame(test_list)

100%|██████████| 3205/3205 [00:12<00:00, 262.88it/s]

```

```
test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3205 entries, 0 to 3204
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   path        3205 non-null   object
1   age         3205 non-null   object
2   sex         3205 non-null   object
3   ImageId     3205 non-null   object
dtypes: object(4)
memory usage: 100.3+ KB

```

```
test.head(5)
```

	path	age	sex	ImageId
0	./siim/dicom-images-test/_/_/ID_af8eb43f4.dcm	22	M	ID_af8eb43f4
1	./siim/dicom-images-test/_/_/ID_114609bb6.dcm	19	M	ID_114609bb6
2	./siim/dicom-images-test/_/_/ID_2e2105281.dcm	41	M	ID_2e2105281
3	./siim/dicom-images-test/_/_/ID_e7eac7dab.dcm	41	M	ID_e7eac7dab
4	./siim/dicom-images-test/_/_/ID_7aef3400f.dcm	83	F	ID_7aef3400f

```
print("Columns having null values in train", train.columns[train.isnull().any()].tolist())
```

```
Columns having null values in train []
```

```
print("No of images having no mask", train[train['lenOfEncodedPixels'] == 0].shape[0])
```

```
No of images having no mask 42
```

```
train.loc[train.lenOfEncodedPixels == 0 , 'encodedPixels'] = np.array(-1)
```

2. Data Preparation

```
def computeMasks(train, path):
```

```

mask_paths = []
for index, row in tqdm(train.iterrows()):
    pixels = row['encodedPixels']
    mask = np.zeros((1024, 1024))
    if row['has_pneumothorax'] == 1:
        for pix in pixels:
            mask = mask + rle2mask(pix, 1024, 1024).T
    img = Image.fromarray(mask).convert('L')
    path2save = path + str(index+1) + ".jpg"
    img.save(path2save)
    mask_paths.append(path2save)
train["mask"] = mask_paths
return train

```

```

path = "siim/dicom-mask-train/"
try:
    os.makedirs(path)
except:
    pass
train = computeMasks(train, path)
train.head(2)

```

12089it [06:14, 32.28it/s]

	path	age	sex	ImageId	enco
0	./siim/dicom-images-train/1.2.276.0.7230010.3....	57	F	1.2.276.0.7230010.3.1.4.8323329.13904.15178752...	
1	./siim/dicom-images-train/1.2.276.0.7230010.3....	52	F	1.2.276.0.7230010.3.1.4.8323329.5535.151787518...	

```

def convertImagesToJpeg(train, path):
    images_paths = []
    for index, row in tqdm(train.iterrows()):
        path2dicom = row['path']
        ds = pydicom.read_file(path2dicom) # read dicom image
        img = ds.pixel_array # get image array
        img_mem = Image.fromarray(img) # Creates an image memory from an object exporting the
        path2save = path + str(index + 1) + ".jpg"
        img_mem.save(path2save)
        images_paths.append(path2save)
    train["images_paths"] = images_paths
    return train

```

```

path = "siim/dicom-images-train-jpg/"
try:
    os.makedirs(path)
except:

```



```
pass
train = convertImagesToJpeg(train, path)
train.head(2)
```

12089it [06:00, 33.51it/s]

	path	age	sex	ImageId	enco
0	./siim/dicom-images-train/1.2.276.0.7230010.3....	57	F	1.2.276.0.7230010.3.1.4.8323329.13904.15178752...	
1	./siim/dicom-images-train/1.2.276.0.7230010.3....	52	F	1.2.276.0.7230010.3.1.4.8323329.5535.151787518...	

```
path = "siim/dicom-images-test-jpg/"
try:
    os.makedirs(path)
except:
    pass
test = convertImagesToJpeg(test, path)
test.head(2)
```

3205it [01:32, 34.79it/s]

	path	age	sex	ImageId	images_paths
0	./siim/dicom-images-test/_/_/ID_bf9328240.dcm	38	F	ID_bf9328240	siim/dicom-images-test-jpg/1.jpg
1	./siim/dicom-images-	52	M	ID_a102h23fd	siim/dicom-images-test-

3. Data Pipeline

```
import tensorflow as tf
# tf.compat.v1.enable_eager_execution()
from tensorflow import keras
from tensorflow.keras.layers import *
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Model, load_model
from tensorflow.keras.layers import UpSampling2D
from tensorflow.keras.layers import MaxPooling2D, GlobalAveragePooling2D
from tensorflow.keras.layers import concatenate
from tensorflow.keras.layers import Multiply
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras import backend as K
from tensorflow.keras.layers import Input, Add, Dense, Activation, ZeroPadding2D, BatchNormal
from tensorflow.keras.models import Model, load_model
from tensorflow.keras.utils import plot_model
from tensorflow.keras.initializers import glorot_uniform, he_normal
from tensorflow.keras.losses import binary_crossentropy
```

```

from tensorflow.keras.datasets import binary_crossentropy
K.set_image_data_format('channels_last')

fileNames = train['images_paths']
labels = train['has_pneumothorax']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(fileNames, labels, test_size=0.10, strati

X_train.head(2)

      8489      siim/dicom-images-train-jpg/8490.jpg
      9227      siim/dicom-images-train-jpg/9228.jpg
      Name: images_paths, dtype: object

y_train.head(2)

      8489      1
      9227      0
      Name: has_pneumothorax, dtype: int64

img_size = 256

def read_image(datapoint):
    image = tf.keras.preprocessing.image.load_img(datapoint)
    image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
    image = tf.image.resize(image, [img_size, img_size])
    image = image / 255.0
    image = exposure.equalize_adapthist(image)      # contrast correction
    return image

def read_label(datapoint):
    datapoint = tf.reshape(datapoint, [1])
    return datapoint

def preprocess(image, label):
    def f(image, label):

        image = image.decode()
        image = read_image(image)

        a = np.random.uniform()
        if a>=0.80:
            b = np.random.uniform()
            if b<0.50:
                image = tf.image.flip_left_right(image)
            else:
                image = tf.image.flip_up_down(image)




```

```
        label = read_label(label)
        return image, label
    images, labels = tf.numpy_function(f, [image, label], [tf.float32, tf.int64])
    images.set_shape([img_size, img_size, 3])
    labels.set_shape([1])
    return images, labels

def tf_dataset(x, y, batch=8):
    dataset = tf.data.Dataset.from_tensor_slices((x, y))
    dataset = dataset.shuffle(buffer_size=15000)
    dataset = dataset.map(preprocess)
    dataset = dataset.batch(batch)
    dataset = dataset.prefetch(2)
    return dataset

train_data = tf_dataset(X_train.values, y_train.values)

for image, result in train_data.take(1):
    sample_image, sample_result = image, result
for i in range(4):
    plt.figure(figsize=(6, 6))
    plt.title("Image")
    plt.imshow(tf.keras.preprocessing.image.array_to_img(sample_image[i]))
    plt.show()
    plt.close()
    print(sample_result[i])
```





```
def read_image(datapoint):
    image = tf.keras.preprocessing.image.load_img(datapoint)
    image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
    image = tf.image.resize(image, [img_size, img_size])
    image = image / 255.0
    image = exposure.equalize_adapthist(image)    # contrast correction
    return image

def read_label(datapoint):
    datapoint = tf.reshape(datapoint, [1])
    return datapoint

def preprocess(image, label):
    def f(image, label):

        image = image.decode()
        image = read_image(image)

        label = read_label(label)
        return image, label
    images, labels = tf.numpy_function(f, [image, label], [tf.float32, tf.int64])
    images.set_shape([img_size, img_size, 3])
    labels.set_shape([1])
    return images, labels

def tf_dataset(x, y, batch=8):
    dataset = tf.data.Dataset.from_tensor_slices((x, y))
    dataset = dataset.shuffle(buffer_size=15000)
    dataset = dataset.map(preprocess)
    dataset = dataset.batch(batch)
    dataset = dataset.prefetch(2)
    return dataset

test_data = tf_dataset(X_test.values, y_test.values)
```

6. Model

```
from sklearn import metrics
```

```

class BasicBlock(tf.keras.layers.Layer):

    def __init__(self, filters, stride=1):
        super(BasicBlock, self).__init__()
        self.conv1_basic = Conv2D(filters=filters, kernel_size=(3, 3), strides=stride, padding
        self.batchNorm1_basic = BatchNormalization()
        self.conv2_basic = Conv2D(filters=filters, kernel_size=(3, 3), strides=1, padding="same"
        self.batchNorm2_basic = BatchNormalization()
        if stride != 1:
            self.downsample_basic = tf.keras.Sequential()
            self.downsample_basic.add(Conv2D(filters=filters, kernel_size=(1, 1), strides=str
            self.downsample_basic.add(BatchNormalization())
        else:
            self.downsample_basic = lambda x: x

    def call(self, inputs, training=None, **kwargs):
        residual = self.downsample_basic(inputs)

        x = self.conv1_basic(inputs)
        x = self.batchNorm1_basic(x, training=training)
        x = tf.nn.relu(x)
        x = self.conv2_basic(x)
        x = self.batchNorm2_basic(x, training=training)

        output = tf.nn.relu(tf.keras.layers.add([residual, x]))

        return output

def make_basic_block_layer(filter_num, blocks, stride=1):
    res_block = tf.keras.Sequential()
    res_block.add(BasicBlock(filter_num, stride=stride))

    for _ in range(1, blocks):
        res_block.add(BasicBlock(filter_num, stride=1))

    return res_block

class Resnet34(tf.keras.Model):
    def __init__(self, input_shape, layer_params):
        super().__init__()

        self.conv1_resnet = tf.keras.layers.Conv2D(filters=32,
                                                    kernel_size=(7, 7),
                                                    strides=2,
                                                    padding="same")
        self.bn1_resnet = tf.keras.layers.BatchNormalization()
        self.pool1_resnet = tf.keras.layers.MaxPool2D(pool_size=(3, 3),
                                                    strides=2,
                                                    padding="same")

```

```

self.layer1_resnet = make_basic_block_layer(filter_num=32,
                                             blocks=layer_params[0])
self.layer2_resnet = make_basic_block_layer(filter_num=64,
                                             blocks=layer_params[1],
                                             stride=2)
self.layer3_resnet = make_basic_block_layer(filter_num=128,
                                             blocks=layer_params[2],
                                             stride=2)
self.layer4_resnet = make_basic_block_layer(filter_num=256,
                                             blocks=layer_params[3],
                                             stride=2)

self.avgpool_resnet = tf.keras.layers.AveragePooling2D()
self.flatten = tf.keras.layers.Flatten()

self.dense1 = tf.keras.layers.Dense(300, activation = 'relu')
self.dense2 = tf.keras.layers.Dense(50, activation = 'relu')
self.final = tf.keras.layers.Dense(1, activation= 'sigmoid')

```

```

def call(self, inputs, training=None, mask=None):
    x = self.conv1_resnet(inputs)
    x = self.bn1_resnet(x, training=training)
    x = tf.nn.relu(x)
    x = self.pool1_resnet(x)
    x = self.layer1_resnet(x, training=training)
    x = self.layer2_resnet(x, training=training)
    x = self.layer3_resnet(x, training=training)
    x = self.layer4_resnet(x, training=training)
    x = self.avgpool_resnet(x)
    x = self.flatten(x)
    x = self.dense1(x)
    x = self.dense2(x)
    output = self.final(x)
    return output

```

```
model3 = Resnet34((256, 256, 3),[3, 4, 6, 3])
```

```
class History_Callback(tf.keras.callbacks.Callback):
```

```

    def on_epoch_end(self, epoch, logs={}):

        val_predict = []
        for i in range(len(X_test)):
            image = tf.keras.preprocessing.image.load_img(X_test.iloc[i])
            image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
            image = tf.image.resize(image, [img_size, img_size])
            image = image / 255.0
            image = exposure.equalize_adapthist(image)

```

```

image = tf.expand_dims(image, axis = 0)
predict = self.model.predict(image)
val_predict.append(np.squeeze(predict, axis = -1).round())

val_targ = y_test
f1_score = metrics.f1_score(val_targ, val_predict)
auc = metrics.roc_auc_score(val_targ, val_predict)
recall = metrics.recall_score(val_targ, val_predict)

print("recall: {recall} - f1Score: {f1score} - AUC: {AUC}".format(f1score = f1_score

history_Model1 = History_Callback()

callbacks = [
    tf.keras.callbacks.ModelCheckpoint('./best_model3.h5', save_weights_only=True, save_best_
        mode='min'),
        history_Model1
]
optimizer = tf.keras.optimizers.Adam(0.001)

model3.compile(optimizer, loss = 'binary_crossentropy' , metrics=['accuracy'])

history3 = model3.fit(train_data, steps_per_epoch=len(train_data), epochs=25,\
        validation_data=test_data,callbacks=callbacks, )

1360/1360 [=====] - 613s 451ms/step - loss: 0.5160 - accuracy: 0.5
recall: 0.0 - f1Score: 0.0 - AUC: 0.5
Epoch 7/25
1360/1360 [=====] - 642s 472ms/step - loss: 0.5149 - accuracy: 0.5
recall: 0.0 - f1Score: 0.0 - AUC: 0.5
Epoch 8/25
1360/1360 [=====] - 640s 471ms/step - loss: 0.5180 - accuracy: 0.5
recall: 0.0 - f1Score: 0.0 - AUC: 0.5
Epoch 9/25
1360/1360 [=====] - 648s 477ms/step - loss: 0.4952 - accuracy: 0.5543
recall: 0.5543071161048689 - f1Score: 0.33789954337899547 - AUC: 0.5324614136787614
Epoch 10/25
1360/1360 [=====] - 652s 479ms/step - loss: 0.4936 - accuracy: 0.0449
recall: 0.0449438202247191 - f1Score: 0.07894736842105263 - AUC: 0.5092022710465421
Epoch 11/25
1360/1360 [=====] - 643s 473ms/step - loss: 0.4767 - accuracy: 1.0
recall: 1.0 - f1Score: 0.3617886178861789 - AUC: 0.5
Epoch 12/25
1360/1360 [=====] - 640s 471ms/step - loss: 0.4812 - accuracy: 0.0
recall: 0.0 - f1Score: 0.0 - AUC: 0.4994692144373673
Epoch 13/25
1360/1360 [=====] - 641s 472ms/step - loss: 0.4668 - accuracy: 0.0823
recall: 0.08239700374531835 - f1Score: 0.1423948220064725 - AUC: 0.5305827906200052
Epoch 14/25
1360/1360 [=====] - 647s 476ms/step - loss: 0.4446 - accuracy: 0.6292
recall: 0.6292134831460674 - f1Score: 0.4132841328413284 - AUC: 0.6139697988978745

```



```

Epoch 15/25
1360/1360 [=====] - 643s 473ms/step - loss: 0.4382 - accuracy: 0.0 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5
Epoch 16/25
1360/1360 [=====] - 646s 475ms/step - loss: 0.4338 - accuracy: 0.11610486891385768 - recall: 0.1962025316455696 - f1Score: 0.1962025316455696 - AUC: 0.5484982943295403
Epoch 17/25
1360/1360 [=====] - 647s 476ms/step - loss: 0.4111 - accuracy: 0.19101123595505617 - recall: 0.2982456140350877 - f1Score: 0.2982456140350877 - AUC: 0.5827667644743434
Epoch 18/25
1360/1360 [=====] - 639s 470ms/step - loss: 0.4198 - accuracy: 0.25842696629213485 - recall: 0.375 - f1Score: 0.375 - AUC: 0.6122283451418211
Epoch 19/25
1360/1360 [=====] - 647s 476ms/step - loss: 0.4114 - accuracy: 0.449438202247191 - recall: 0.45197740112994345 - f1Score: 0.45197740112994345 - AUC: 0.6482859801044873
Epoch 20/25
1360/1360 [=====] - 643s 473ms/step - loss: 0.3909 - accuracy: 0.1797752808988764 - recall: 0.27507163323782235 - f1Score: 0.27507163323782235 - AUC: 0.5718409313199265
Epoch 21/25
1360/1360 [=====] - 644s 474ms/step - loss: 0.3787 - accuracy: 0.00749063670411985 - recall: 0.014705882352941176 - f1Score: 0.014705882352941176 - AUC: 0.502152961664161
Epoch 22/25
1360/1360 [=====] - 663s 488ms/step - loss: 0.3872 - accuracy: 0.0449438202247191 - recall: 0.08362369337979093 - f1Score: 0.08362369337979093 - AUC: 0.518225625611298
Epoch 23/25
1360/1360 [=====] - 650s 478ms/step - loss: 0.3794 - accuracy: 0.4606741573033708 - recall: 0.5061728395061729 - f1Score: 0.5061728395061729 - AUC: 0.6793816646389466
Epoch 24/25
1360/1360 [=====] - 655s 482ms/step - loss: 0.3641 - accuracy: 0.18352059925093633 - recall: 0.28654970760233917 - f1Score: 0.28654970760233917 - AUC: 0.5779598749970181
Epoch 25/25
1360/1360 [=====] - 664s 488ms/step - loss: 0.3548 - accuracy:

```

#this is after changed pipeline

```

history3 = model3.fit(train_data, steps_per_epoch=len(train_data), epochs=1,\
                      validation_data=test_data,callbacks=callbacks, )

```

2/1360 [.....] - ETA: 55:39 - loss: 0.6746 - accuracy: 0.875

```
history3 = model3.fit(train_data, steps_per_epoch=len(train_data), epochs=25,\
                      validation_data=test_data,callbacks=callbacks, )
```

Epoch 6/25

1360/1360 [=====] - 635s 467ms/step - loss: 0.5231 - accuracy: 0.5231 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5

Epoch 7/25

1360/1360 [=====] - 633s 465ms/step - loss: 0.5227 - accuracy: 0.5227 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5

Epoch 8/25

1360/1360 [=====] - 632s 464ms/step - loss: 0.5216 - accuracy: 0.5216 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5

Epoch 9/25

1360/1360 [=====] - 639s 470ms/step - loss: 0.5199 - accuracy: 0.5199 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5

Epoch 10/25

1360/1360 [=====] - 635s 467ms/step - loss: 0.5210 - accuracy: 0.5210 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5

Epoch 11/25

1360/1360 [=====] - 644s 473ms/step - loss: 0.5161 - accuracy: 0.5161 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5

Epoch 12/25

1360/1360 [=====] - 635s 467ms/step - loss: 0.5096 - accuracy: 0.5096 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5

Epoch 13/25

1360/1360 [=====] - 636s 467ms/step - loss: 0.5082 - accuracy: 0.5082 - recall: 0.003745318352059925 - f1Score: 0.007434944237918215 - AUC: 0.50134187361339

Epoch 14/25

1360/1360 [=====] - 636s 467ms/step - loss: 0.5039 - accuracy: 0.5039 - recall: 0.0 - f1Score: 0.0 - AUC: 0.4989384288747346

Epoch 15/25

1360/1360 [=====] - 637s 469ms/step - loss: 0.4923 - accuracy: 0.4923 - recall: 0.0 - f1Score: 0.0 - AUC: 0.5

Epoch 16/25

1360/1360 [=====] - 637s 468ms/step - loss: 0.4799 - accuracy: 0.4799 - recall: 0.14606741573033707 - f1Score: 0.2154696132596685 - AUC: 0.5433097163577375

Epoch 17/25

1360/1360 [=====] - 641s 472ms/step - loss: 0.4911 - accuracy: 0.4911 - recall: 0.026217228464419477 - f1Score: 0.04778156996587031 - AUC: 0.503023688542188

Epoch 18/25

1360/1360 [=====] - 642s 472ms/step - loss: 0.4784 - accuracy: 0.4784 - recall: 0.0599250936329588 - f1Score: 0.1111111111111111 - AUC: 0.5273086190033158

Epoch 19/25

1360/1360 [=====] - 653s 480ms/step - loss: 0.4664 - accuracy: 0.4664 - recall: 0.026217228464419477 - f1Score: 0.04982206405693951 - AUC: 0.509393115293781

Epoch 20/25

1360/1360 [=====] - 650s 478ms/step - loss: 0.4580 - accuracy: 0.4580 - recall: 0.27715355805243447 - f1Score: 0.35576923076923084 - AUC: 0.5987678618287651

Epoch 21/25

1360/1360 [=====] - 654s 481ms/step - loss: 0.4521 - accuracy: 0.4521 - recall: 0.033707865168539325 - f1Score: 0.06405693950177936 - AUC: 0.514200004771106

Epoch 22/25

1360/1360 [=====] - 650s 478ms/step - loss: 0.4495 - accuracy: 0.4495 - recall: 0.2247191011235955 - f1Score: 0.3225806451612903 - AUC: 0.5884742002433265

Epoch 23/25

1360/1360 [=====] - 654s 481ms/step - loss: 0.4376 - accuracy: 0.1947565543071161 - f1Score: 0.29714285714285715 - AUC: 0.5809239247119444

Epoch 24/25

1360/1360 [=====] - 663s 488ms/step - loss: 0.4259 - accuracy: 0.4794007490636704 - f1Score: 0.48669201520912553 - AUC: 0.670167465826952

Epoch 25/25

```
history3 = model3.fit(train_data, steps_per_epoch=len(train_data), epochs=20,\n                      validation_data=test_data,callbacks=callbacks, )
```

Epoch 1/20

1360/1360 [=====] - 681s 500ms/step - loss: 0.4165 - accuracy: 0.38202247191011235 - f1Score: 0.4646924829157176 - AUC: 0.6538562465707675

Epoch 2/20

1360/1360 [=====] - 686s 504ms/step - loss: 0.4053 - accuracy: 0.2958801498127341 - f1Score: 0.3940149625935162 - AUC: 0.6187468689615688

Epoch 3/20

1360/1360 [=====] - 700s 515ms/step - loss: 0.4017 - accuracy: 0.5243445692883895 - f1Score: 0.5303030303030303 - AUC: 0.6979472315656384

Epoch 4/20

1360/1360 [=====] - 691s 508ms/step - loss: 0.3959 - accuracy: 0.3408239700374532 - f1Score: 0.44067796610169496 - AUC: 0.6412187790739283

Epoch 5/20

1360/1360 [=====] - 697s 512ms/step - loss: 0.3892 - accuracy: 0.4307116104868914 - f1Score: 0.5066079295154186 - AUC: 0.6771392447338915

Epoch 6/20

1360/1360 [=====] - 695s 511ms/step - loss: 0.3911 - accuracy: 0.49063670411985016 - f1Score: 0.5229540918163673 - AUC: 0.6906474391087575

Epoch 7/20

1360/1360 [=====] - 713s 524ms/step - loss: 0.3845 - accuracy: 0.25842696629213485 - f1Score: 0.3822714681440444 - AUC: 0.6159438440802499

Epoch 8/20

1360/1360 [=====] - 692s 509ms/step - loss: 0.3826 - accuracy: 0.4157303370786517 - f1Score: 0.49443207126948785 - AUC: 0.6701793935924043

Epoch 9/20

1360/1360 [=====] - 691s 508ms/step - loss: 0.3745 - accuracy: 0.3595505617977528 - f1Score: 0.4393592677345538 - AUC: 0.6404971492640569

Epoch 10/20

1360/1360 [=====] - 687s 505ms/step - loss: 0.3644 - accuracy: 0.4344569288389513 - f1Score: 0.5178571428571428 - AUC: 0.6827274028483503

Epoch 11/20

1360/1360 [=====] - 679s 499ms/step - loss: 0.3602 - accuracy: 0.3857677902621723 - f1Score: 0.4963855421686747 - AUC: 0.6689985448126149

Epoch 12/20

1360/1360 [=====] - 673s 495ms/step - loss: 0.3516 - accuracy: 0.42696629213483145 - f1Score: 0.5241379310344828 - AUC: 0.6848207256852501

Epoch 13/20

1360/1360 [=====] - 637s 468ms/step - loss: 0.3426 - accuracy: 0.5543071161048689 - f1Score: 0.5077186963979418 - AUC: 0.6879815835301415

Epoch 14/20

1360/1360 [=====] - 642s 472ms/step - loss: 0.3465 - accuracy: 0.3857677902621723 - f1Score: 0.4768518518518518 - AUC: 0.659975190247859

Epoch 15/20

1360/1360 [=====] - 666s 490ms/step - loss: 0.3211 - accuracy: 0.3970037453183521 - f1Score: 0.4988235294117647 - AUC: 0.6709010234022759

```

Epoch 16/20
1360/1360 [=====] - 653s 480ms/step - loss: 0.3081 - accuracy: 0.4794
recall: 0.4794007490636704 - f1Score: 0.5267489711934157 - AUC: 0.6913988883322598
Epoch 17/20
1360/1360 [=====] - 660s 485ms/step - loss: 0.2942 - accuracy: 0.4831
recall: 0.48314606741573035 - f1Score: 0.516 - AUC: 0.6863713351940648
Epoch 18/20
1360/1360 [=====] - 647s 476ms/step - loss: 0.2873 - accuracy: 0.5056
recall: 0.5056179775280899 - f1Score: 0.531496062992126 - AUC: 0.6965457191249792
Epoch 19/20
1360/1360 [=====] - 657s 483ms/step - loss: 0.2747 - accuracy: 0.3558
recall: 0.35580524344569286 - f1Score: 0.46798029556650245 - AUC: 0.6545480569670078
Epoch 20/20
1360/1360 [=====] - 657s 483ms/step - loss: 0.2747 - accuracy: 0.3558
recall: 0.35580524344569286 - f1Score: 0.46798029556650245 - AUC: 0.6545480569670078

```

```
model3.summary()
```

```
Model: "resnet34_1"
```

Layer (type)	Output Shape	Param #
=====		
conv2d_36 (Conv2D)	multiple	4736
batch_normalization_36 (Batch Normalization)	multiple	128
max_pooling2d_1 (MaxPooling2D)	multiple	0
sequential_7 (Sequential)	(None, 64, 64, 32)	56256
sequential_8 (Sequential)	(None, 32, 32, 64)	281408
sequential_10 (Sequential)	(None, 16, 16, 128)	1712256
sequential_12 (Sequential)	(None, 8, 8, 256)	3285760
average_pooling2d_1 (Average Pooling2D)	multiple	0
flatten_1 (Flatten)	multiple	0
dense_3 (Dense)	multiple	1229100
dense_4 (Dense)	multiple	15050
dense_5 (Dense)	multiple	51
=====		
Total params: 6,584,745		
Trainable params: 6,576,233		
Non-trainable params: 8,512		

```
model3.load_weights("best_model3.h5")
```

```
#Saving the model
```

```
model3.save("pneumothorax_classifier")
```

```
WARNING:absl:Found untraced functions such as conv2d_37_layer_call_and_return_conditio
WARNING:absl:Found untraced functions such as conv2d_37_layer_call_and_return_conditio
INFO:tensorflow:Assets written to: pneumothorax_classifier/assets
INFO:tensorflow:Assets written to: pneumothorax_classifier/assets
```

```
!zip -r /content/pneumothorax_classifier.zip /content/pneumothorax_classifier/
```



```
adding: content/pneumothorax_classifier/ (stored 0%)
adding: content/pneumothorax_classifier/variables/ (stored 0%)
adding: content/pneumothorax_classifier/variables/variables.index (deflated 80%)
adding: content/pneumothorax_classifier/variables/variables.data-00000-of-00001 (defla
adding: content/pneumothorax_classifier/assets/ (stored 0%)
adding: content/pneumothorax_classifier/saved_model.pb (deflated 92%)
```

```
# Plot training & validation iou_score values
```

```
plt.figure(figsize=(30, 5))
plt.subplot(121)
plt.plot(history3.history['loss'])
plt.plot(history3.history['val_loss'])
plt.title('Model loss')
plt.ylabel('loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
```

```
# Plot training & validation loss values
```

```
plt.subplot(122)
plt.plot(history3.history['accuracy'])
plt.plot(history3.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



-- Oversampling the images which are predicted wrongly

```
train.iloc[0]['images_paths']

'siim/dicom-images-train-jpg/1.jpg'

imagesIndexes = []
for i in tqdm(range(len(train))):
    image = tf.keras.preprocessing.image.load_img(train.iloc[i]['images_paths'])
    image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
    image = tf.image.resize(image, [img_size, img_size])
    image = image / 255.0
    image = exposure.equalize_adapthist(image)
    image = tf.expand_dims(image, axis = 0)
    prediction = model3.predict(image)

    if np.squeeze(prediction, -1).round() != train.iloc[i]['has_pneumothorax']:
        if prediction <= 0.3 or prediction >= 0.7:
            imagesIndexes.append(i)
```

```
100%|██████████| 12089/12089 [25:01<00:00, 8.05it/s]
```

```
len(imagesIndexes)
```

```
910
```

```
fileNames = train['images_paths']
labels = train['has_pneumothorax']
```

```
fileNames.shape
```

```
(12089,)
```

```
#Copy pasting so that these images appear two times
wrongPredictionImages = pd.Series(train.iloc[imagesIndexes]['images_paths'])
wrongPredictionLabels = pd.Series(train.iloc[imagesIndexes]['has_pneumothorax'])
```

```
fileNames = fileNames.append(wrongPredictionImages, ignore_index=True)
labels = labels.append(wrongPredictionLabels, ignore_index= True)
```

```
fileNames.shape
```

```
(12999,)
```

```
fileNames.tail(2)
```

```
12997     siim/dicom-images-train-jpg/12053.jpg
12998     siim/dicom-images-train-jpg/12073.jpg
Name: images_paths, dtype: object
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(fileNames, labels, test_size=0.10, strati
```

```
def read_image(datapoint):
    image = tf.keras.preprocessing.image.load_img(datapoint)
    image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
    image = tf.image.resize(image, [img_size, img_size])
    image = image / 255.0
    image = exposure.equalize_adapthist(image)      # contrast correction
    return image

def read_label(datapoint):
    datapoint = tf.reshape(datapoint, [1])
    return datapoint

def preprocess(image, label):
    def f(image, label):

        image = image.decode()
        image = read_image(image)

        a = np.random.uniform()
        if a>=0.80:
            b = np.random.uniform()
            if b<0.50:
                image = tf.image.flip_left_right(image)
            else:
                image = tf.image.flip_up_down(image)

        label = read_label(label)
        return image, label
    images, labels = tf.numpy_function(f, [image, label], [tf.float32, tf.int64])
    images.set_shape([img_size, img_size, 3])
    labels.set_shape([1])
    return images, labels

def tf_dataset(x, y, batch=8):
    dataset = tf.data.Dataset.from_tensor_slices((x, y))
    dataset = dataset.shuffle(buffer_size=15000)
    dataset = dataset.map(preprocess)
    dataset = dataset.batch(batch)
    dataset = dataset.prefetch(2)
    return dataset
```

```
train_data = tf_dataset(X_train.values, y_train.values)
```

```

def read_image(datapoint):
    image = tf.keras.preprocessing.image.load_img(datapoint)
    image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
    image = tf.image.resize(image, [img_size, img_size])
    image = image / 255.0
    image = exposure.equalize_adapthist(image)      # contrast correction
    return image

def read_label(datapoint):
    datapoint = tf.reshape(datapoint, [1])
    return datapoint

def preprocess(image, label):
    def f(image, label):
        image = image.decode()
        image = read_image(image)

        label = read_label(label)
        return image, label
    images, labels = tf.numpy_function(f, [image, label], [tf.float32, tf.int64])
    images.set_shape([img_size, img_size, 3])
    labels.set_shape([1])
    return images, labels

def tf_dataset(x, y, batch=8):
    dataset = tf.data.Dataset.from_tensor_slices((x, y))
    dataset = dataset.shuffle(buffer_size=15000)
    dataset = dataset.map(preprocess)
    dataset = dataset.batch(batch)
    dataset = dataset.prefetch(2)
    return dataset

test_data = tf_dataset(X_test.values, y_test.values)

model4 = Resnet34((256, 256, 3),[3, 4, 6, 3])

callbacks = [
    tf.keras.callbacks.ModelCheckpoint('./best_model4.h5', save_weights_only=True, save_best_
                                     mode='min'),
    history_Model1
]
optimizer = tf.keras.optimizers.Adam(0.001)

model4.compile(optimizer, loss = 'binary_crossentropy' , metrics=['accuracy'])

#this is after changed pipeline

```



```
history4 = model4.fit(train_data, steps_per_epoch=len(train_data), epochs=1,\
                      validation_data=test_data,callbacks=callbacks, )
```

2/1360 [.....] - ETA: 52:28 - loss: 2.5245 - accuracy: 0.506

KeyboardInterrupt Traceback (most recent call last)

<ipython-input-46-403a4323ca77> in <module>()

1 #this is after changed pipeline

```
----> 2 history4 = model4.fit(train_data, steps_per_epoch=len(train_data), epochs=1,\
validation_data=test_data,callbacks=callbacks, )
```

⌵ 6 frames

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/execute.py in

quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)

58 ctx.ensure_initialized()

59 tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,

```
---> 60 inputs, attrs, num_outputs)
```

61 except core._NotOkStatusException as e:

62 if name is not None:

KeyboardInterrupt:

SEARCH STACK OVERFLOW

```
model4.load_weights("best_model3.h5")
```

```
history4 = model4.fit(train_data, steps_per_epoch=len(train_data), epochs=4,\
                      validation_data=test_data,callbacks=callbacks, )
```

Epoch 1/4

1463/1463 [=====] - 717s 490ms/step - loss: 0.3934 - accuracy: recall: 0.5644699140401146 - f1Score: 0.6283891547049442 - AUC: 0.7396482062314137

Epoch 2/4

1463/1463 [=====] - 717s 490ms/step - loss: 0.3730 - accuracy: recall: 0.6819484240687679 - f1Score: 0.6770981507823612 - AUC: 0.7799857788062031

Epoch 3/4

1463/1463 [=====] - 725s 496ms/step - loss: 0.3619 - accuracy: recall: 0.4899713467048711 - f1Score: 0.5927209705372616 - AUC: 0.7150172190937605

Epoch 4/4

1463/1463 [=====] - 721s 493ms/step - loss: 0.3482 - accuracy: recall: 0.47564469914040114 - f1Score: 0.5724137931034482 - AUC: 0.7036477964682026

```
# Plot training & validation iou_score values
```

```
plt.figure(figsize=(30, 5))
```

```
plt.subplot(121)
```

```
plt.plot(history4.history['loss'])
```

```
plt.plot(history4.history['val_loss'])
```



```
plt.title('Model loss')
```

```
plt.ylabel('loss')
```

```
plt.xlabel('Epoch')
```

```
plt.legend(['Train', 'Test'], loc='upper left')
```

```
# Plot training & validation loss values
plt.subplot(122)
plt.plot(history4.history['accuracy'])
plt.plot(history4.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



```
imagesIndexesOverSample = []
for i in tqdm(range(len(train))):
    image = tf.keras.preprocessing.image.load_img(train.iloc[i]['images_paths'])
    image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
    image = tf.image.resize(image, [img_size, img_size])
    image = image / 255.0
    image = exposure.equalize_adapthist(image)
    image = tf.expand_dims(image, axis = 0)
    prediction = model4.predict(image)

    if np.squeeze(prediction, -1).round() != train.iloc[i]['has_pneumothorax']:
        if prediction <= 0.3 or prediction >= 0.7:
            imagesIndexesOverSample.append(i)
```

100%|██████████| 12089/12089 [25:05<00:00, 8.03it/s]

```
len(imagesIndexesOverSample)
```

484

```
#Saving the model
model4.save("pneumothorax_classifier_oversampling")
```

WARNING:absl:Found untraced functions such as conv2d_37_layer_call_and_return_conditio
 WARNING:absl:Found untraced functions such as conv2d_37_layer_call_and_return_conditio

```
INFO:tensorflow:Assets written to: pneumothorax_classifier_oversampling/assets
INFO:tensorflow:Assets written to: pneumothorax_classifier_oversampling/assets
```

```
!zip -r /content/pneumothorax_classifier_oversampling.zip /content/pneumothorax_classifier_ov

adding: content/pneumothorax_classifier_oversampling/ (stored 0%)
adding: content/pneumothorax_classifier_oversampling/variables/ (stored 0%)
adding: content/pneumothorax_classifier_oversampling/variables/variables.index (deflated 0%)
adding: content/pneumothorax_classifier_oversampling/variables/variables.data-000001-of-000001 (deflated 0%)
adding: content/pneumothorax_classifier_oversampling/assets/ (stored 0%)
adding: content/pneumothorax_classifier_oversampling/saved_model.pb (deflated 91%)
```

Training on highly missclassified points reduced our highly missclassified points to half.

Comparison of two models

```
from random import sample
listOfIndices = sample(imagesIndexes, 2)
listOfIndices

[7566, 9532]

for index in listOfIndices:
    image = tf.keras.preprocessing.image.load_img(train.iloc[index]['images_paths'])
    image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
    image = tf.image.resize(image, [img_size, img_size])
    image = image / 255.0
    image = exposure.equalize_adapthist(image)
    image = tf.expand_dims(image, axis = 0)
    #Old model
    prediction = model3.predict(image)
    print("Prediction using old model - (before oversampling) Actual Value :", train.iloc[index]['Actual Value'])
    #New model
    prediction = model4.predict(image)
    print("Prediction using new model - (after oversampling) Actual Value :", train.iloc[index]['Actual Value'])
    print("\n")

Prediction using old model - (before oversampling) Actual Value : 1 predicted value: [0]
Prediction using new model - (after oversampling) Actual Value : 1 predicted value: [1]

Prediction using old model - (before oversampling) Actual Value : 1 predicted value: [0]
Prediction using new model - (after oversampling) Actual Value : 1 predicted value: [1]
```

- Model is performing better after oversampling

Lets try changing the threshold

```
model4.load_weights("best_model4.h5")
```

```
def roundingOff(value, threshold):
    return 1 if value >=threshold else 0
```

```
def findingBestThreshold(thresholdList):
    truePositives = np.zeros(len(thresholdList))
    falseNegatives = np.zeros(len(thresholdList))
    for i in tqdm(range(len(train))):
        image = tf.keras.preprocessing.image.load_img(train.iloc[i]['images_paths'])
        image = tf.keras.preprocessing.image.img_to_array(image, dtype='float32')
        image = tf.image.resize(image, [img_size, img_size])
        image = image / 255.0
        image = exposure.equalize_adapthist(image)
        image = tf.expand_dims(image, axis = 0)
        prediction = model4.predict(image)

        y_true = train.iloc[i]['has_pneumothorax']
        for thres in range(len(thresholdList)):
            y_pred = roundingOff(prediction, thresholdList[thres])

            if y_true == 1 and y_pred == 1:
                truePositives[thres] = truePositives[thres] + 1
            elif y_true == 1 and y_pred == 0:
                falseNegatives[thres] = falseNegatives[thres] + 1
    print()
    for thres in range(len(thresholdList)):
        print("Recall for threshold ", thresholdList[thres], "is: ", (truePositives[thres])/(
```

```
thresholdList = [0.4, 0.45, 0.5, 0.6, 0.75]
findingBestThreshold(thresholdList)
```

```
100%|██████████| 12089/12089 [31:13<00:00, 6.45it/s]
Recall for threshold 0.4 is: 0.8130385912326714
Recall for threshold 0.45 is: 0.7853128512551517
Recall for threshold 0.5 is: 0.735481453727988
Recall for threshold 0.6 is: 0.6403147246159611
Recall for threshold 0.75 is: 0.4166354439865118
```

```
thresholdList = [0.38, 0.39, 0.4, 0.41, 0.42]
findingBestThreshold(thresholdList)
```

```
100%|██████████| 12089/12089 [31:02<00:00, 6.49it/s]
```

```
Recall for threshold 0.38 is: 0.8291494941925814
Recall for threshold 0.39 is: 0.8227800674409891
Recall for threshold 0.4 is: 0.8130385912326714
Recall for threshold 0.41 is: 0.8077931809666542
Recall for threshold 0.42 is: 0.8029224428624954
```

- 0.38 is the best threshold.
- Further increasing the recall is resulting in decrease in precision so we are stopping at this point.

